
How To Build a ROS Robot

SDSMT Robotics Team

Ian Carlson

August 8, 2015

Contents

Title	i
Contents	i
List of Figures	iii
1 Introduction	1
1.1 Purpose	1
2 Mechanical Design	3
2.1 Design Considerations	3
2.2 Requirements	4
2.3 Kinematic Models	4
2.3.1 Differential Drive	4
2.3.2 Ackermann Drive	5
2.3.3 Skid Steer	6
2.3.4 Omni-Driv	7
3 Electrical Design	9
3.1 Design Considerations	9
3.1.1 The Basics	9
3.1.2 Voltage	9
3.1.3 Current	10
3.1.4 Resistance	10
3.1.5 Blinking Lights	11
3.1.6 Voltage Divider	13
3.2 Batteries	14
3.2.1 The Basics	14
3.2.2 Lead Acid Batteries - Pb	14
3.2.3 Lithium Polymer - LiPo	14
3.2.4 Other Batteries	15
3.3 Tips For Not Lighting Things On Fire	15
4 Odroid Setup	17
4.1 Purpose Of This Chapter	17
4.2 Before You Begin	17
4.3 Obtain the Fuser	17
4.4 Run The Fuser	18
4.5 Loading The EMMC Image	18
Bibliography	19

List of Figures

2.1	Differential Drive	5
2.2	Ackermann Drive	6
2.3	Skid Steer	6
2.4	SMP Robot	7
2.5	Mecanum Robot	8
2.6	Mecanum Wheel	8
3.1	Resistors In Series	10
3.2	Resistors In Parallel	11
3.3	Parallel Resistors Example	11
3.4	Parallel Resistors Example Reduced	11
3.5	LED Proper Hookup	12
3.6	LED Proper Hookup - With Voltage Drop	12
3.7	Differential Drive	13
3.8	LiPo Voltage VS Charge	15

1

Introduction

1.1 Purpose

Building a sophisticated robot can be a complex and potentially daunting task. While a single afternoon can be sufficient to strap a few servos and an arduino to a chassis, doing anything interesting required an investment of both time and resources. As the cost of a robot increases, it becomes more important to have a solid plan to follow. This document attempts to outline the major phases of robotics development, as well as providing the basic information required to build a basic robot featuring ROS. It should be noted that this document is entirely a product of my own experiences - mostly the failures - during my time on the Robotics team, in the hope that others will not repeat them. That being said, this document is neither expected to be entirely complete nor perfectly accurate. It is a certain truth that failure is a better teacher than success, but experience is the best teacher of all - preferably someone else's.

The major sections are as follows:

- Mechanical Design - Basic mobile robot kinematics and design principles.
- Electrical Design - How not start your robot on fire. Probably.
- Setting Up The Odroid - How to go from a fresh Odroid to one running Fedora 22 and ROS
- ROS Architecture - The basic design for a remote controlled robot in ROS

This guide should be delivered alongside its companion document - Robot Paper Template - and potentially the example I created using that template - Mecanum and SMP Design Document. All of the code referenced in this document is available on the SDSMT Robotics Team GitHub pages - SMD-ROS-DEVEL and SMD-ROBOTICS.

2

Mechanical Design

2.1 Design Considerations

The most important aspect of your robot is ensuring that it is mechanically capable of fulfilling requirements - and making sure it doesn't collapse under its own weight. No amount of sophisticated code, nor the most powerful computer in the world, can force your robot to function if the chassis is compromised. It is tempting - we fell in to this trap several times during my time on the Robotics Team - to want to design your perfect robot from the ground up. Don't. If this is your first robot, or even second or third, take my word for it and just buy a chassis.

Seriously.

Mechanical design is time-consuming and, unless you have several skilled mechanical designers around, likely going to be difficult. On the Mines campus, it has become increasingly difficult to do any real fabrication, thanks in part to the increasingly strict attention to safety, and the CAT lab's understandable desire to cut down on machining costs.

Going for it anyway? Good for you, however, keep a few things in mind:

First, keep it simple. Complex designs are more expensive, have tighter tolerances, take a great deal longer to design and fabricate, and - most importantly - have more ways to fail. The best design in the world is useless if it never makes it off paper.

Second, keep moving parts to a minimum. This is similar to point one, but it's important enough to say twice. Every moving part means you need to deal with routing power i.e. keeping cables from getting tangled. Every moving piece also generally adds weight, cost, and complexity.

Third, if the words "we'll attach this here somehow" are ever uttered, take a closer look. We like to think in the CSC department that Mechanical Engineering is just hitting things with hammers and turning wrenches, but building real things is actually very difficult. Every piece of your robot has to be attached somehow - avoid glue - and that's not something that just sorts itself out once you have all the pieces.

Fourth - No, hot glue is not a good idea. Trust me. Been there. It's not. Epoxy can work, but it's no substitute for bolts.

Fifth - 3M velcro tape is better, but still not good, and is only appropriate for an indoor non-competition research vehicle where frequent reconfiguration is expected. See the SMP and Mecanum chassis for examples. It doesn't hold up under vibration as well as you would think. It holds great against normal stress, but not terribly well against sheer stress, and not at all against twisting stress.

Sixth, iterate. Don't expect your first build to be perfect. In all honesty, it probably won't even work. You forgot something or overlooked something or assumed something that isn't right. I promise. Give me a call before you start designing, and I will bet you \$100. Expect your design to fail, and get something built as early as possible. A failed design isn't the end of the world unless it's the day of competition. More than one of our designs failed because we were half-way into the spring semester before we got anything built, and there was no time to try again.

Seventh, failure is good. You learn a whole lot from mistakes. Don't be afraid to try something just because it might not work. It *definitely* won't work the first time anyway.

2.2 Requirements

What does your robot need to do? Agile development is great for software, and there are some applications that can be adapted to hardware, but you really need to know what you're getting into before you start buying or building things. I can't possibly predict every possible requirement, but here are a few important ones I have come across before.

How does the robot need to move? Whether mobile or fixed, your robot likely needs to actuate within its workspace, and this will be one of the main factors in the mechanical design. This will be addressed further in section 2.3.

How fast does the robot need to be? If your robot doesn't need to be very fast, you can reduce power requirements, use smaller and/or cheaper motors, and you don't have to worry as much about vibrations. If the robot needs to be fast, the inverse is true.

How strong does the robot need to be? If the robot only needs to drive from A to B on a card table carrying a few paperclips, an Arduino taped to a few continuous rotation servos and some cardboard will probably do the trick. If you need to carry a cinderblock around an outdoor obstacle course, you need a robust frame, bigger motors, bigger batteries, and space to put it all.

How durable does the robot need to be? This one is pretty self-explanatory. However, keep in mind that as soon as you have a working RC robot, people will want to play with it. They will crash it. Hard. Once the robot is driving autonomously, it will crash itself. Often *and* hard. Make sure all potential contact surfaces on the robot can survive impact at maximum speed with something sharp and/or hard.

Does it need to be waterproof or water-resistant? Some competitions - like the Sparkfun AVC - are outdoors. Given Murphy's law, this means it will probably rain. Plan early about ways to keep water out and the magic smoke in. Never put holes in the top, and consider rubber grommets or seals around connectors and cabling. Also consider having angled non-porous surfaces on the top of the robot to deflect water such that it will drip harmlessly away from the electronics.

Does it need to be heat-tolerant? This, unfortunately, is in direct conflict with being waterproof, because you need to put holes in the robot for air circulation. We never came up with a perfect solution that handles both. Assume your CPU, GPU (if you have one), motor controller, and motors will attempt to melt anything in contact with them as soon as you close the case. Make sure you have proper heat-sinking and, if possible, forced air-flow over those heat-sinks. Also make sure the air has a way to escape. If operating outdoors, make sure you have filtered air intake.

Finally, what does it need to do? Claws, grippers, baskets, sensors, and payloads all need secure mounting. Make sure you include some sort of mounting bracket for whatever tools your robot will need to complete its job.

2.3 Kinematic Models

Now that you know what your robot needs to do, it's time to actually start designing something. For a mobile robot, kinematics are everything. There are a few common kinematic models that I want to cover in some detail, including the math. It is very important to understand - at least conceptually - the mathematics behind robot motion if you want to have any hope for autonomy.

Each kinematic model has different advantages and disadvantages, so it is important to consider how your robot needs to be able to move before selecting one.

2.3.1 Differential Drive

The simplest and most common drive system for a mobile robot is differential drive. This involves a fixed axle with two independently driven motors at the ends. A picture is worth a thousand words, so check out Figure 3.1.6, borrowed from the nice people at www.robotplatform.com.

Differential Drive is very simple because it only requires two powered wheels. In the image, the front wheel is a caster wheel, or a skid-pad, and is only present to keep the robot upright.

One major advantage of the differential drive paradigm is the ability to turn in place by spinning one wheel forward, and one back. However, because of the skid-pad or caster wheel, differential drive is typically

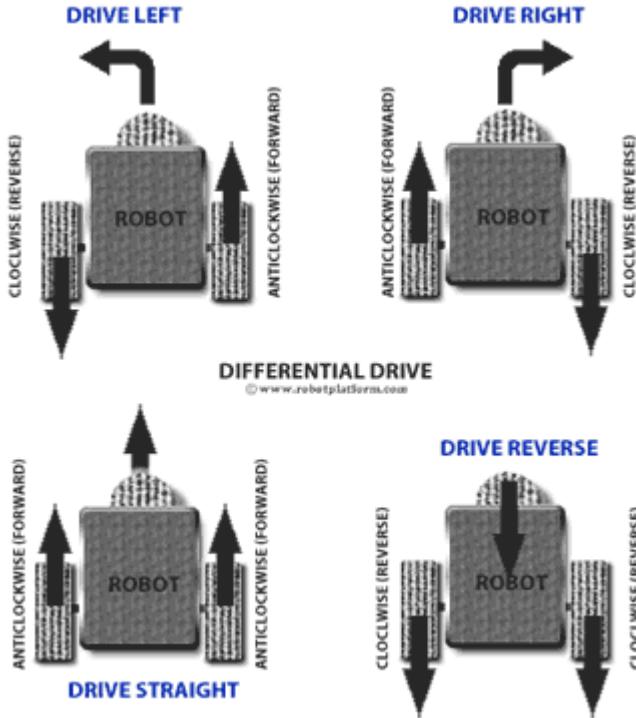


Figure 2.1: Differential Drive

only used in low-speed applications. The extra drag and wobbling of an uncontrolled wheel could cause vibrations or instability at high speeds.

I won't go into the derivation of the kinematic model, but I will present it here so you can use it for reference. There are two kinematic models of interest, the forward kinematic model and the reverse.

The reverse kinematic model tells you how the robot will move given a set of wheel velocities, and is shown in Equation ??.

TODO: Add kinematics?

2.3.2 Ackermann Drive

Ackermann drive is a common choice for high speed robots because it provides better stability and less drag than the differential drive. It is also the same steering system that most cars use. However, an Ackermann steered vehicle cannot turn in place.

An Ackermann drive vehicle must have two front wheels which can be turned, and two fixed rear wheels. It is common to have powered rear wheels since the pivoting front wheels adds additional complexity when attempting to mount motors, but powered front wheels are not unheard of. See Figure ??.

There is one additional complication to consider when attempting to build your own Ackermann Drive system. When an Ackermann vehicle turns, each wheel is essentially drawing out the circumference of a circle. To take that thought a step further, both steering wheels will be drawing out circles, but of slightly different diameters - the difference being the center to center distance from wheel to wheel. Therefore, to avoid wheel slip while drawing two circles of different sizes, each wheel must be at a slightly different angle. The difference changes depending on how sharp of a turn is desired. Obviously designing a mechanical system to automatically handle this would be complex, and writing code to handle turning each wheel individually would be non-trivial at best. I will once again recommend buying a professionally made drive system over building one unless you have a really good reason not to.

TODO: Add kinematics?

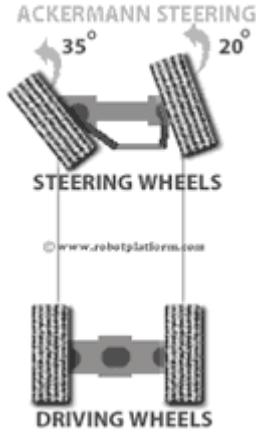


Figure 2.2: Ackermann Drive

2.3.3 Skid Steer

Skid Steer A.K.A. Tank Drive, is another common form of locomotion. It's similar to having two differential drive systems mounted with a rigid bar between them. See Figure 2.3.3 for an example of how Skid Steer drives. Each side, regardless of how many wheels it has, is unified, meaning that all wheels on the left drive at the same time and speed, as do all the wheels on the right. Tank drive is arguably better for an outdoor robot because it is much harder to end up high-centered when all of your wheels are powered. This also gives you more traction. However, to do anything other than drive straight forward or backward, some of the wheels have to slip. The largest disadvantage to Skid Steer is that any sort of wheel speed or encoder based odometry becomes very unreliable.

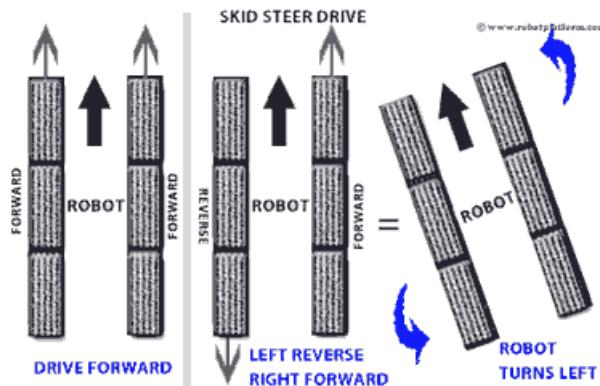


Figure 2.3: Skid Steer

However, if you have some other means of localization, like the LIDAR on the SMP robots, or the ASUS RGBD sensor, this becomes less of a problem. You end up with less drag than traditional differential drive and potentially more power because each wheel may (but doesn't have to) have its own motor. In addition, this design is much simpler to implement than Ackermann Steering or Omni-Drive, which will be discussed next. For these reasons, the SMP robots use Skid Steer. See the SMP in Figure 2.3.3.

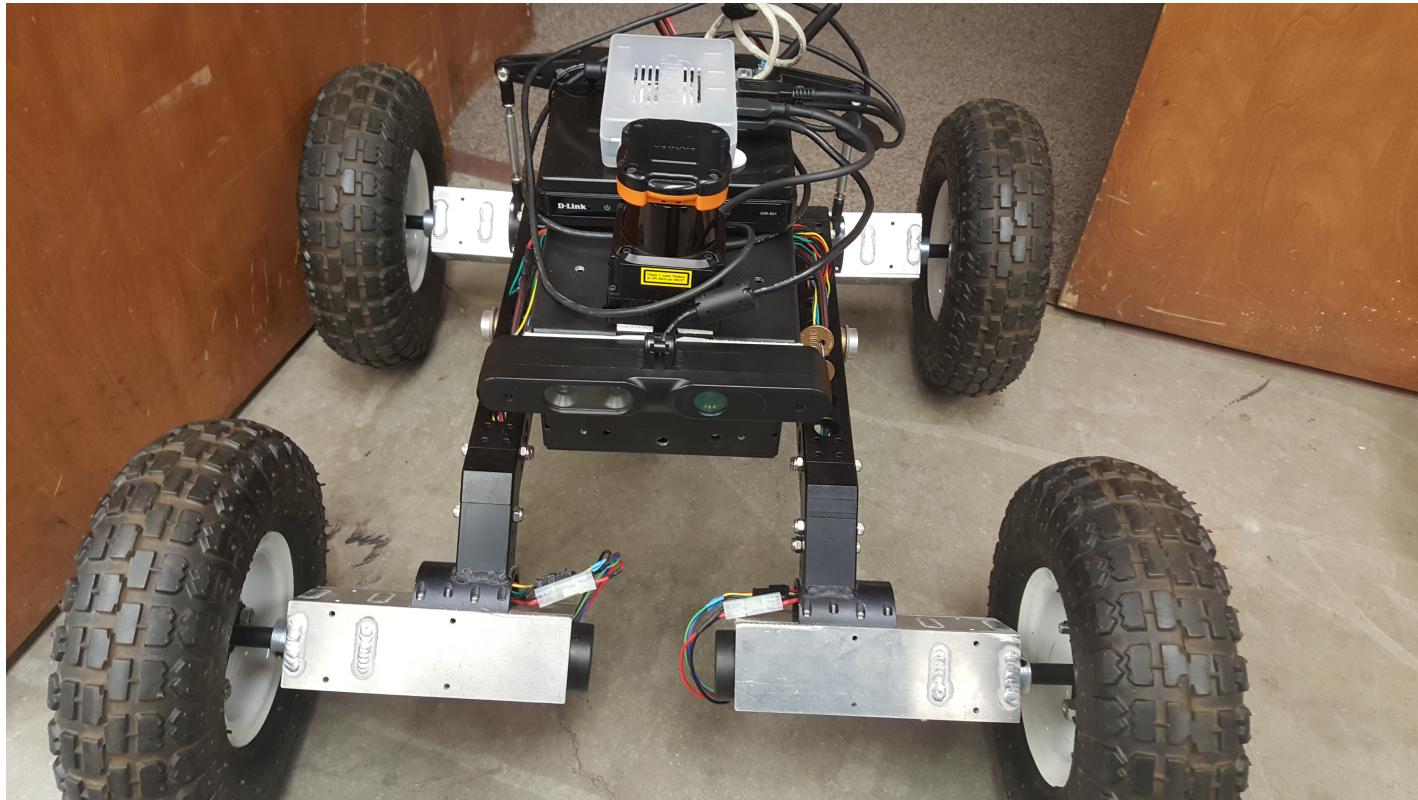


Figure 2.4: SMP Robot

2.3.4 Omni-Driv

An Omni-Drive robot is one that can move in any direction. There are several ways to construct an Omni-Drive robot, but I will focus on a four wheeled robot with fixed wheels in "traditional" positions. The Mecanum robot shown in Figure ?? - named so for the style of wheels, shown in Figure 2.3.4 - is one such robot.

Omni-Drive kinematics are complex, and out of the scope for this guide. However, I will attempt a brief summary. The rollers on a mecanum wheel are mounted at 45 degrees. This means that of any amount of force applied to rotating a wheel, only approximately half will be exerted along the line parallel with the direction of travel. The other half will be exerted along the perpendicular axis. The wheels are not mounted symmetrically, however. All wheels are mounted such that, if all wheels turn forward, the perpendicular forces will point inward and cancel, allowing the robot to move forward. By operating the wheels in opposing pairs, a force in any direction can be generated, allowing the robot to translated and rotate in any direction freely. This is very useful for a robot, because it means there aren't any constraints that have to be considered when attempting to operate the robot autonomously. If you want to move in a certain way, you can, no questions asked.

While the actual mathematics for driving an Omni-Drive robot are complex, the concept is simple. Just calculate the translation you want to perform, then the rotation, and add them. In this case, the total is just the sum of the parts. Special care should be taken, however. Attempting to both drive forward and turn at full speed works fine mathematically, but once you command a wheel to drive faster than its top speed, the robot will begin to behave erratically, as the forces will no longer properly cancel.

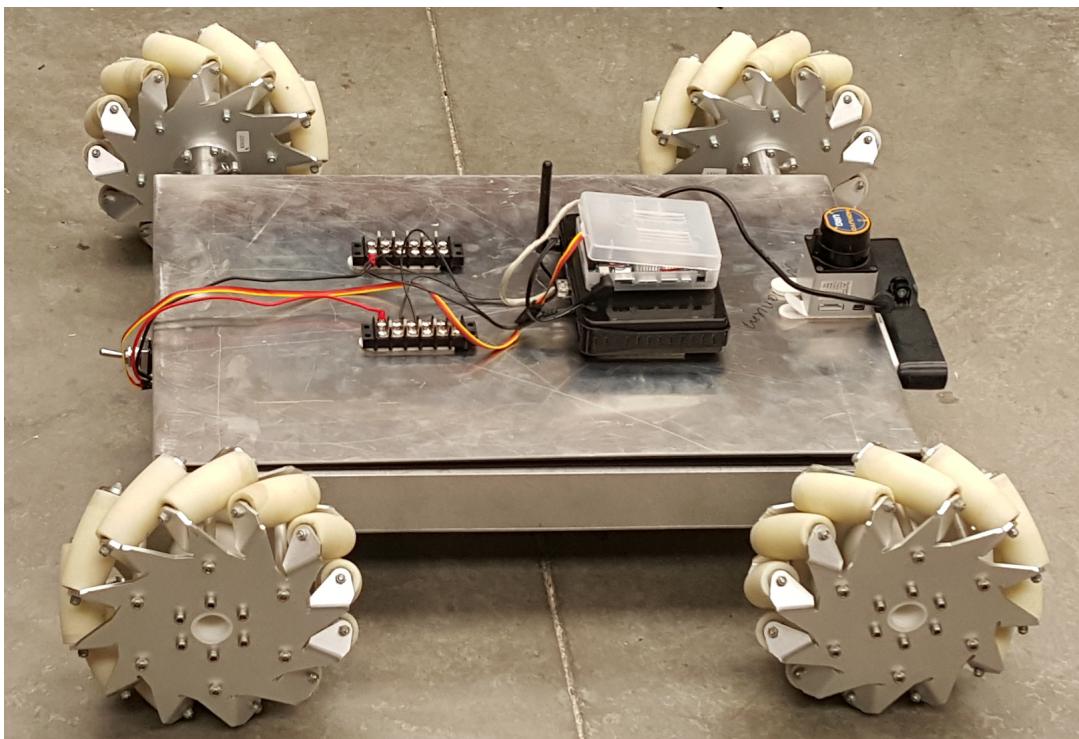


Figure 2.5: Mecanum Robot



Figure 2.6: Mecanum Wheel

3

Electrical Design

3.1 Design Considerations

The most important aspect of your robot is ensuring that it is electrically capable of fulfilling requirements - and making sure it doesn't light itself on fire or electrocute anyone. The best mechanical design in the world is useless if there's nothing to supply power to the wheels, and your computer is just a hunk of plastic and metal if it doesn't get the right amount of power. While electrical design for a robot is not very complicated (unless you want it to be), it does require some fundamental knowledge.

3.1.1 The Basics

We'll start off easy. I'm going to restrict the discussion here to (relatively) low voltage DC circuits. Within that domain, there are three very important terms. Voltage, measured in Volts (V), Current, measured in Amperes "Amps" (A), and Resistance, measured in Ohms Ω . A very smart German dude by the name of Georg Ohm came up with probably the most important (for us at least) law in all of electronics back in 1827, and published it in his book *Die galvanische Kette, mathematisch bearbeitet*. (Translated to English, that's *The Galvanic Circuit Investigated Mathematically*). The book is available online, but I wouldn't recommend it unless you're in for some heavy reading. It was written in the early 1800s after all. The important part - the law we care about here - is called Ohm's law and is shown in Equation 3.1.

$$V = IR \tag{3.1}$$

That's it. That's all there is to it. Voltage (V) equals Current (I) times Resistance(R). This law has the interesting property that, unlike everything else we do in Electrical Engineering, it almost always holds true. So, if something isn't behaving the way you think it should, refer to this law.

Ok, so now we have the fundamental law of the universe, but so what? What are all these funky terms? Let's talk more about that. Sparkfun also has a pretty excellent page explaining this, so if my explanation doesn't work out for you, just search for *Sparkfun Voltage, Current, Resistance, and Ohm's Law*.

3.1.2 Voltage

Voltage behaves very much like water pressure in a pipe. With water, pressure - or more accurately, pressure differential - is what causes water to flow. However, just because there is pressure doesn't mean the water actually flows. It's just a measure of how much the water *wants* to flow. For example, consider a completely full, sealed container of water. Now shrink the volume of the container by half without letting any water out. Some laws that I vaguely remember from chemistry say that the pressure has to increase. (Or temperature, but likely it will be some of both). Now there's more pressure, but the container is still sealed with no route for the water to escape, so there's no flow.

Voltage behaves almost exactly the same. Except instead of water, we have a high density of electrons. Because they are negatively charged, and like charges repel one another, electrons like to spread out as much as they can to reach the minimum energy configuration. Sort of like when you drop a glass of water on the

floor, and the water spreads out to get as low to the ground as possible - thereby reducing it's gravitational potential energy and reaching a minimal energy state. So, compressing a lot of electrons in to a small space creates this electrical pressure - voltage, or electromotive force for people who like words with many syllables.

The common source of voltage differentials in mobile robotics are batteries or DC power supplies.

3.1.3 Current

If voltage is pressure, than current is the amount of flow. If there is a pressure differential, and a route for water to escape, then flow happens in a pipe. This flow continues until the pressure differential is balanced out. Batteries work the same way. (By the way, think of power supplies as infinite batteries.) A fully charged battery has some voltage, and if there is a path from the negative terminal to the positive terminal, electrons will flow along that path. As the electrons flow from negative to positive and thereby "spread out" to the minimal energy state, the battery voltage will decrease. This flow is current.

3.1.4 Resistance

Resistance is not a property of the electricity itself, but a property of the material it flows through. The easy analogy is to think of the resistance as the size of the pipe. A narrower pipe allows less water to flow through. Likewise, a material with higher resistances "resists" the flow of electrons more strongly. In low power DC applications, you can generally assume that wires have a resistance of 0.

Often, we will use components called "Resistors" to purposefully introduce current in a circuit. This is particularly useful when we have components like LEDs (light-emitting diodes) which can only handle so much current. We'll have an example circuit a little later demonstrating how to choose the right resistor for an LED.

There is one more thing you need to know about resistors before we move on. What happens if you have more than one? There are two ways to hook up multiple resistors. You can put them in parallel or in series. In Figure 3.1.4 - from wikipedia, which covers this subject in great depth- there are N resistors hooked up in series. This case is very easy. The value of the resistors just adds up, so you can pretend like you have one big resistor with a value that is the sum of all the resistors in series. Because I like math, the oh-so-difficult math for this is shown in Equation 3.2.

$$R = R_1 + R_2 + \dots + R_n \quad (3.2)$$

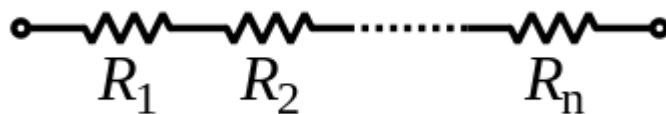


Figure 3.1: Resistors In Series

The other way to hook up resistors is in parallel. This one is a little more complicated. An example is shown in Figure 3.1.4.

Putting It All Together

Example time! First, look at the picture in Figure ???. In that example, there are multiple paths for the electricity to follow. Meditating on Ohm's law for a while will eventually lead you to the mathematical reason behind this, but for now just take my word for it: the electricity divides itself up among the possible paths inversely proportional to the resistance of each path. Ohm's law requires this, and it's actually convenient for us because it generates the least possible waste heat - but that's an advanced topic I won't cover here. Let's attach some real numbers to this example. Consider a circuit with a 10V voltage source and 2 resistors

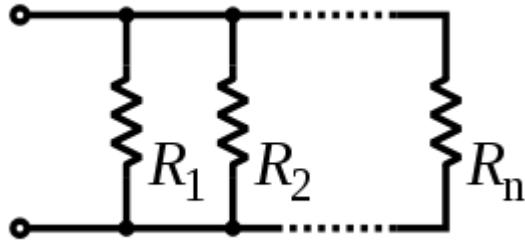


Figure 3.2: Resistors In Parallel

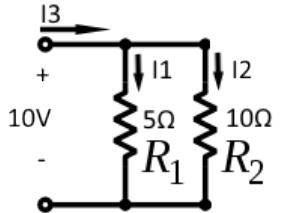


Figure 3.3: Parallel Resistors Example

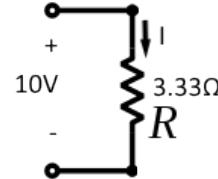


Figure 3.4: Parallel Resistors Example Reduced

in parallel, one with 5Ω and one with 10Ω . This example circuit is shown in Figure 3.1.4. Reducing the two resistors with Equations ??, gives the circuit showin in Figure 3.1.4.

There are a few things to think about in this example. First, note that I_3 in Figure 3.1.4 must be equivalent to the current I in Figure 3.1.4 since they are in truth the same circuit. In Figure3.1.4, we can easily solve for I using Ohm's law.

$$\begin{aligned}
 V &= IR \\
 10 &= I(3.33) \\
 I &= 10/3.33 \\
 I &= 3A
 \end{aligned} \tag{3.3}$$

We can likewise solve for I_1 and I_2 in Figure 3.1.4 using the same manner, but with R values of 5Ω and 10Ω respectively. This gives us $I_1 = 2A$ and $I_2 = 1A$. Notice that $I_3 = I_1 + I_2$. Another side-effect of Ohm's law is that the sum of all currents entering and leaving a node is 0, which is intuitive if you think about it. It's the same thing as splitting a hose. The sum of the volumes of water that come out of the two hoses must be the same as the volume that goes in to the un-split end.

3.1.5 Blinking Lights

One of the first things most people want to do with circuits is make lights blink. This is good. In the micro-controller world, you often don't have a terminal output until after you get through a decent amount of setup, so blinking LEDs can be the only way to debug problems during boot. The first thing most people do when they want to blink a light is blow one up. This should generally be avoided, although LEDs are cheap and it's a good learning experience.

First, let's look at the proper way to hook up an LED, shown in Figure 3.1.5. The main thing to note here is the resistor. Supplying too much current to an LED lets out the magic smoke, and remember, if you let out the magic smoke, the component no longer works. The next question, is how do you choose a value

for R ? This depends on the LED, and the voltage source. For kicks, let's assume you've got a 5V power supply, from an arduino GPIO for example. By the way, drawing too much power from an arduino GPIO is a good way to kill it, although the LED usually goes first. I'll assume you've got an LED that wants 10mA of current, although this may vary. Check the datasheet for your specific part.

There is one caveat about LEDs. We haven't talked about diodes because it's out of the scope of this document. However, LEDs are a special type of component called a "semi-conductor" which means it is only conductive some of the time. In this case, the LED is only conductive when it has at least a certain threshold of voltage pushing in the "correct" direction. Pay attention to which side the Anode and Cathode of your LED are plugged in to. If it's in the wrong way, current will not flow. (Unless you supply enough voltage to cause the diode to enter the "breakdown" voltage region. In this case, that's the amount of voltage where the LED breaks, but the breakdown region is actually very important for some components. Google "Zener Diode" for more on that. Hint: They're used in motor driver boards for back-emf protection.)

So why do we care about the LED being a semi-conductor? Because semi-conductors usually have a property called a "voltage-drop". This property causes the component to eat a portion of the voltage in the circuit. Your LED datasheet should say something about the voltage drop of the LED, but a typical value is 0.7V for hobby level LEDs. We typically model this voltage drop as a small voltage source pointed in the "wrong" direction. See Figure 3.1.5.

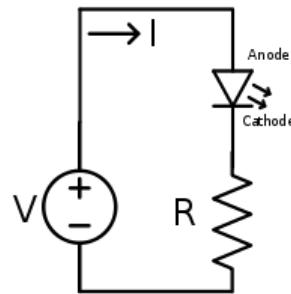


Figure 3.5: LED Proper Hookup

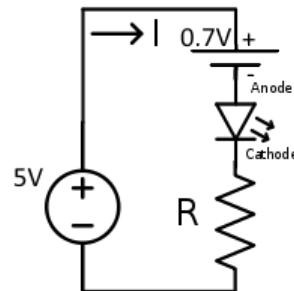


Figure 3.6: LED Proper Hookup - With Voltage Drop

In this case we can just sum the two voltage sources, and pretend we have a 4.3V power supply. So, we have a 4.3V source, and we want a 10mA current. Easy, right? Just use Ohm's law as shown in Equation 3.5.

$$\begin{aligned}
 V &= IR \\
 4.3 &= (.01)R \\
 R &= 4.3/.01 \\
 R &= 430\Omega
 \end{aligned} \tag{3.4}$$

It so happens that LEDs are generally pretty robust. You can probably go plus or minus about 50% on

that R value. A smaller value will make the LED brighter, while a larger value will make it dimmer. Too large a value will burn it out. An excessively large value *may* make it explode. I highly recommend trying it once under controlled circumstances. Wear eye protection! There will probably be exactly one or two pieces of shrapnel, but it would really suck to get one in your eye. For best results, connect the LED to a power supply and slowly turn up the voltage. You'll get to see all phases of operation for the LED. Below 0.7V it probably will not turn on. From 0.7 to somewhere around 4-6V it will steadily get brighter. At some point it will reach a maximum brightness and then begin to heat up, at which point it will probably start to dim. At some point after that, it will go out. Cranking up the voltage more will (probably) cause it to explode. Sometimes the conductive element inside just melts and you don't get an explosion. It really depends on the LED.

3.1.6 Voltage Divider

There's one last circuit I want to mention because it becomes very important when dealing with analog sensors. The voltage divider. An example is shown in Figure ???. The basic idea here is that you have a known resistor value for R_1 , and some variable value for R_2 . A common one is a thermistor - a resistor that changes its value based on temperature. In fact, an excellent experiment to try is to use a thermistor in exactly this circuit configuration to measure the temperature of water. It's very easy, and will teach you about the analog to digital converter. There's also the matter of sensor calibration, but I won't go in to the details on that here. Suffice it to say that all sensors are terrible, and you should never trust them until you've tested them rigorously.

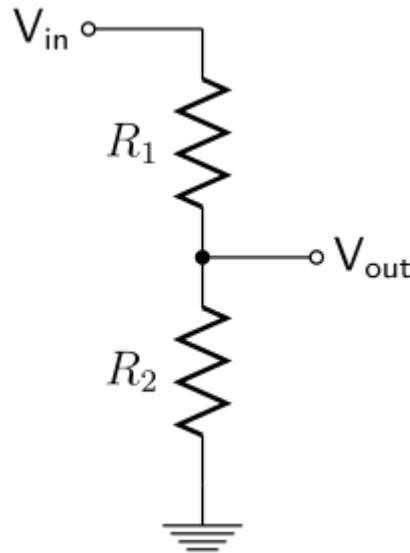


Figure 3.7: Differential Drive

The general plan is to supply a known voltage to V_{in} , have a known R_1 , and measure V_{out} . With these pieces of information, you can solve for R_2 using Ohm's law. The derivation isn't difficult, but other people have already done it, so here it is. Once you have R_2 you just refer to your sensor calibration or the datasheet for the part, and you have your value for (in this case) temperature.

$$R_2 = R_1 * \frac{1}{\left(\frac{V_{in}}{V_{out}} - 1\right)} \quad (3.5)$$

3.2 Batteries

3.2.1 The Basics

We have had a lot of problems dealing with batteries over the years. Partly because people are lazy, but mostly because they don't understand how batteries work or how to properly maintain them. All batteries are basically slow, controlled, chemical reactions. They are not magical devices that store and dispense electric charge. The closest thing to that is a capacitor, which is two metal plates with an insulator between them. By applying voltage, you charged the plates, storing energy in the form of actual electric charge. However, capacitors tend to discharge all of their stored energy at once. In addition, the total energy they can store is far less than most batteries. Batteries, on the other hand, store energy in the form of chemical potential energy. This is far more stable than storing raw electric charge, but it does lead to a few problems. The big one is that, since the energy storage relies on chemistry, temperature is important. Being stored in a place that is too hot or too cold can cause a battery to burst, or drain it. Another is that, over time, the chemistry can get disrupted, causing the battery's maximum storage potential to decline.

3.2.2 Lead Acid Batteries - Pb

Lead acid batteries are very stable. That's why we use them in cars. They also have a pretty good capacity, and are able to source a tremendous amount of current at once. This is good, because it takes a lot of force to turn over an engine. For a rugged, outdoor robot that may experience a variety of temperature conditions, lead-acid may be the way to go. However, lead-acid batteries are generally larger and heavier than their counterparts. Charging them is easy. Just hook them up to a power supply, set the power supply to the battery's rated voltage, and limit the current. What you limit the current to depends on the battery. Car batteries can generally handle up to 6A. The real problem is heat. If you charge too fast, the metal plates in the battery heat up, and can cause the acid to boil. This creates potentially toxic vapors and, if the vapors escape the battery housing, reduce the lifespan and charge of the battery. Other than charging too quickly, there isn't much to worry about here. Please charge in a well-ventilated area, just in case. Running a lead-acid battery dead isn't really a big deal as long as you don't leave it dead for a long time, or it doesn't get to cold while dead.

3.2.3 Lithium Polymer - LiPo

In the robots I built during my time on the team, we used these most commonly. They are light-weight, small, and can store a great deal of power. A LiPo battery generally consists of some number of cells. Each cell has a rated voltage of 3.7V. Batteries with more voltage are built by putting multiple cells in series. This voltage has to do with the internal chemistry. LiPo batteries are not nearly as stable as lead acid. There are three major concerns when dealing with a LiPo. First, only charge using a LiPo charger. There are some extra pins on a LiPo battery that tell the charger important information about the cells within. Second, make sure you look up the rating for charging a LiPo. The general rule of thumb is that a LiPo can be charged at the rate of 1C. If you have a 1500 mAh battery, you can charge it at 1.5A. Charging it slower is fine. Charging faster can cause the LiPo to heat up, swell, and potentially burst. When the LiPo bursts, the chemicals inside will spontaneously burn, creating fire, pressure, and heat. Seriously bad news. Third, never ever cut or puncture a LiPo. An externally ruptured LiPo is bad news, but one that ruptures internally, or without a good pressure relief, is a thermal grenade.

That all said, LiPo batteries are pretty safe if you follow the guidelines I set out for charging. There are two more things to worry about. 3.7V is the rated voltage for a cell, but when you charge, you generally charge to about 4.2V - the charger will handle this. Never manually overcharge the LiPo. However, unlike a lead-acid battery, letting the charge get too low in a LiPo will permanently damage it. The minimum safe level for a single cell is 3V. Always monitor the voltage of LiPo batteries you are using and ensure they don't drop below this level. The difficult part is, LiPo voltage doesn't drop linearly. Figure ?? shows voltage versus charge for a LiPo. As you can see from the graph, you have to monitor the battery voltage very carefully, because it decreases rapidly once the charge gets low.

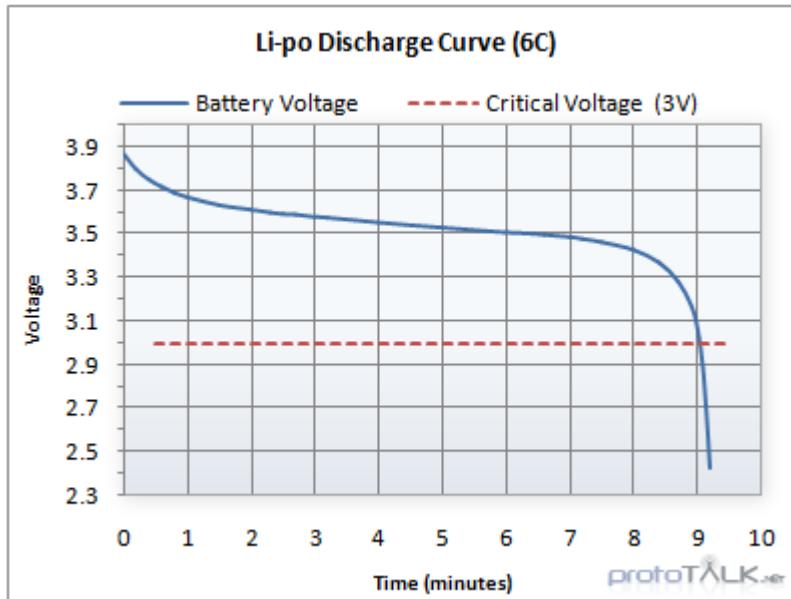


Figure 3.8: LiPo Voltage VS Charge

As a side note, this chart is for one specific battery I found on a forum post at Traxxas.com. Individual results may vary in specifics, but the point is the same. Monitor your voltage and don't let it drop below 3.0V.

The final thing to mention is balancing. Because a LiPo battery may be made up of multiple cells, the total voltage isn't enough to tell the health of the battery. In the lab we have at least two balances. Some chargers will also balance while they charge. Balancing slowly bleeds charge from one cell and puts it into another cell. This keeps the cells from becoming unbalanced. This is a good thing. Each cell may discharge differently because chemistry is a sloppy science and never works quite the way it's supposed to. This can lead to one cell with a much higher or lower voltage than the others. If one cell gets overcharged, it may swell and/or rupture. If one cell gets too low, it may "die". Dead cells are ones that have dropped to a low enough voltage that they cannot safely be recharged. Most chargers will refuse to charge the battery if there are dead cells. If you know what you're doing, sometimes dead cells can be nursed back to life, but it's a delicate and potentially dangerous procedure. It's usually better to recycle the battery and get a new one.

3.2.4 Other Batteries

There are three more common types of batteries. Lithium Iron Phosphate - LiFePO₄ often pronounced "LieFo" - Nickle Metal Hydride - NiMH - and Nickle Cadmium - NiCd "Nigh-Cad". I haven't personally worked with either, so the internet is a better resource than I am. My vague understanding is that LiFePO₄ are similar to LiPo batteries but more stable and more expensive. A quick Google search tells me that NiMH is the new NiCd, and is used in general consumer electronics, so it must be pretty stable, but doesn't (in general) have as much energy density as Lithium batteries.

3.3 Tips For Not Lighting Things On Fire

Forethought and attentiveness are key for keeping the magic smoke inside the components. All electrical components are made using plastic, silicon, some trace metals, and a mystical substance called magic smoke. If you give the component too much voltage, current, or heat the magic smoke will use this extra energy to break free and escape. It is a very clever substance, and even just a momentary spike is enough to free it.

At this point the component will no longer work. Nobody is perfect, but here are a few tips I've picked up over the years to keep the magic smoke locked up tight.

Turn Off The Power

This should be common sense, but you'd be surprised how often people get overconfident about what they're doing and modify a circuit while it's powered. Sure, if you know what you're doing you're theoretically safe. The real problem arises when one of those tiny wires gets away from you. Powered wires appear to be supernaturally attracted to conductive terminals, particularly ones that will cause sparks. Just don't do it. The five seconds it takes to flip off the power could save you three days of waiting for new parts.

Electrical Tape

If you're working on a circuit, if you aren't immediately dealing with a wire, tape the end. Most of the instances in which I blew something up were because I forgot about a wire that wasn't connected, and it brushed up against something else, making a short. This is particularly important when dealing with batteries. Remember, you can't turn a battery off. If both terminals of a battery aren't connected to something, the free wires should be taped over.

Fun Fact: If you plug the gate pin on a MOSFET in to 24VDC, while the source is floating and the drain is grounded, it explodes.

Be Careful Around Batteries

I've alluded to this several times, but I'll say it again. Always be careful around batteries. I had six years of robotics team experience, and four years of that was while taking courses in Electrical Engineering. At the end of my last semester, I was putting different connectors on some LiPo batteries. I thought to myself, "I'll just cut both wires on the battery at the same time so they'll be the same length."

Technically speaking, it worked out. Turns out the reflex when you hear a loud SNAP and see a bright flash of light is to close your fist, which for me finished cutting the wires and yanked the tool away from the battery cables. It could have been worse though. Just think about what you're doing around batteries. Remember, they can't be turned off.

Don't Work When You're Tired

Or drunk. Or sick. Or angry. It makes you inattentive, and then things like the above happen. It doesn't matter how smart you are. Tired people make mistakes.

Ground Loops

Ground loops are tricky devils. Explaining the whole mechanism behind them is outside the scope of this document, however, if you are ever dealing with a circuit with multiple power sources, look them up. A ground loop can ruin your whole day. Essentially, ground loops occur when two things you consider "ground" have some resistance between them. If that happens, current through one loop induces a voltage difference in the other loop. Wikipedia is your friend here. Moral of the story - use isolated power supplies, and be very careful when plugging multiple power sources in to one set of electronics. When in doubt, unplug your computer from the wall before plugging anything grounded into it. USB cables, for instance.

Use the Mult-Meter

The multi-meter is your friend. When powering up a system for the first time, unplug all sensitive electronics, power the system, and check voltages at your main terminals. Make sure the value *and* the polarity are correct. Then power down, plug in just one piece of equipment, and power back up. Repeat until everything is plugged in. This is useful for two reasons. First, if you have a problem, it will be very obvious which piece is the culprit. Second, if something blows, you risk damaging the minimum possible number of expensive items.

4

Odroid Setup

4.1 Purpose Of This Chapter

This chapter is intended to be a guide for setting up an Odroid XU3 with Fedora 21, using an EMMC as the storage device. Step by step guides for using an SD card exist in plenty on the internet, and only the first two steps are any different. The steps should work with Fedora 22 as well, but Fedora 22 was still very new when I last worked with the Odroids and wasn't ready for installation on an arm architecture yet. For installation on other Odroid platforms, the steps should be very similar, but may not be exactly the same.

4.2 Before You Begin

You will need the following items to complete this guide:

1. Odroid XU3
2. USB to MicroSD Adapter
3. MicroSD to EMMC Adapter
4. EMMC
5. Computer Running Fedora 20/21
6. USB to Serial Adapter

Check your version of Fedora before beginning. You can do this using the command:

```
$ cat /etc/issue
```

You want to check if you're running Fedora 20 or Fedora 21+. This will be important later.

4.3 Obtain the Fuser

The Fuser is a utility written by Scott Logan to set up the EMMC in such a way that the Fedora image fits on the EMMC efficiently. It sets the "fuse" bit, which are sort of like software-writable settings on the EMMC. Modifying these bits manually is pretty difficult, which is why you'll want to use the Fuser to do it for you.

On Fedora 21+, run:

```
$ sudo yum install smd-odroid-release  
$ sudo yum install odroid-xu3-sd-fuser
```

These two commands download repository information for the smd-odroid-release repository - which Scott maintains - and then downloads and installs the fuser itself.

Of Fedora 20, the repository headers will not automatically download. You'll have to download the fuser more directly. Instead run:

```
$ sudo yum install http://csc.mcs.sdsmt.edu/smd-odroid/fedora/linux/21/x86_64/odroid-xu3-
```

4.4 Run The Fuser

If you run in to problems, there is a wiki page with instructions specifically for the fuser at <https://github.com/sdsmt-robotics-odroids/odroid-xu/wiki/Boot-Media-Fusing>.

Before inserting the EMMC, run the command:

```
$ ls /dev/mmcblk*
```

Take note of the entries there. Whatever you see there before inserting the EMMC is what you will NOT want to target in the next steps. Next, connect the EMMC to your computer. Plug the EMMC into the EMMC to MicroSD adapter. If your machine has a MicroSD slot, you can plug that directly in to your computer. If not, (or if you encounter problems, which can happen rarely mostly due to driver issues) you'll need a MicroSD to USB adapter.

Run the ls command again. There should be new entries, probably something like /dev/mmcblkX... in a few various forms. Take note of the value of X. Now run the following command, substituting your number (usually 0) for X.

```
$ sudo odroid-xu3-emmc-fuser /dev/mmcblkX
```

You should see the following:

```
Successfully enabled write access to /dev/mmcblk0boot0.
Fusing boot blob to /dev/mmcblk0boot0...
1262+0 records in
1262+0 records out
646144 bytes (646 kB) copied, 1.71663 s, 376 kB/s
Success!
```

If you don't see something like this, specifically if there are 0+0 records in or out, something has gone wrong. Try again or check the wiki instructions for the Fuser.

4.5 Loading The EMMC Image

First you need to get the image to write to the EMMC. The images are hosted at <http://csc.mcs.sdsmt.edu/smd-odroid/fedora/linux/21/Images/armhf/>. You'll want to download the "latest" image for your type of Odroid. In this case, the filename will probably be something like: f21-odroid-xu3-minimal-latest.img.xz.

Once you have downloaded the image, uncompress it using the following command with <foo> replaced by the image name. This will take about a minute.

```
$ xz -d <foo>.img.xz
```

Next, you need to umount the EMMC partitions or they won't write correctly. Do so with the following command, again substituting your value for X:

```
$ sudo umount /dev/mmcblkXp*
```

Be careful with this next command. dd is a very powerful tool, and it can totally erase your harddrive if you let it. Mostly just make sure you don't accidentally point it at the wrong device. if is the input file - the image you uncompressed. of is the target device, the mmcblkX we identified earlier. Substitute your value for X. bs is the block size. Just leave it alone.

```
$ sudo dd if=<foo>.img of=/dev/mmcblkX bs=8M
```

This takes a while. On the order of 3 or 4 minutes when I did it last. Be patient. When it finishes, you should expect to see:

```
204+1 records in
204+1 records out
1713373184 bytes (1.7 GB) copied, 91.7598 s, 18.7 MB/s
```

Next, to make sure the OS has flushed all the buffers to the device, run:

```
$ sync
```

If there is a noticeable pause after running sync, run it again just in case. That pause means it's working on something. If there is no noticeable pause, it didn't do anything, and you're good to go.

4.6 What Next?

At this point, you have a functional EMMC with Fedora 21 installed. Go ahead and plug it back in to the Odroid. I recommend getting a USB-Serial converter cable to debug with. I also recommend getting a CMOS battery for the Odroid, or the system clock will reset every time the board is restarted. This makes yum unhappy and you'll have a bad time. ROS can throw some very strange time-related errors as well. If you only needed Fedora installed, you're done with this chapter - except perhaps see the section on common error messages. Otherwise, move on to the next section to install ROS. It's possible the Scott has the kick-start version up by now, with ROS already installed, but at the time of writing, that hasn't happened yet.

4.7 Installing ROS

The ROS build for arm isn't in with the normal ROS or Linux repositories. It was in fact hosted from, when I left, a set of 10 Odroids over in M113. To install this repository, run:

```
$ sudo yum install http://csc.mcs.sdsmt.edu/smd-pub/fedora/linux/21/armhf/smd-ros-release
```

You'll also need rpm fusion, which is an extended set of packages for Fedora not bundled in to the main set of rpms. To install those, run:

```
$ sudo yum localinstall --nogpgcheck http://download1.rpmfusion.org/free/fedora/rpmfusion-repo-f18.noarch.rpm http://download1.rpmfusion.org/nonfree/fedora/rpmfusion-nonfree-repo-f18.noarch.rpm
```

Now, you should be able to use yum to install pretty much any ROS Fedora package. The following command will install the basic starting ROS packages. It doesn't include PCL or Gazebo, but you shouldn't be running those on Odroids anyway.

```
$ sudo yum install ros-indigo-robot
```

If you want to do on-board compiling, which you almost certainly will unless you've already finished development and are just deploying functional binaries, run the following command:

```
$ sudo yum install python-rosdep python-rosinstall_generator python-wstool python-rosinstall
```

That's it. Installing ROS is way easier than it used to be. If in doubt, here's a link to the ROS wiki page dealing with installation issues: <http://wiki.ros.org/indigo/Installation/Source>.

This is a fairly basic install, so there may be some unresolved dependencies. If you're having dependency issues, run the following command where <src> is a directory containing ROS packages.

```
$ rosdep install --from-path <src> --ignore-src
```

4.8 Common Errors

There are a handful of pretty common errors that I found solutions for:

One of the configured repositories failed (Fedora 21 – x86_64 – Updates), and yum doesn't have enough cached data to **continue**. At this point the only safe thing yum can **do** is fail. There are a few ways to work "fix" this:

1. Contact the upstream **for** the repository and get them to fix the problem.
2. Reconfigure the baseurl/etc. **for** the repository, to point to a working upstream. This is most often useful **if** you are using a newer distribution release than is supported by the repository (and the packages **for** the previous distribution release still work).
3. Disable the repository, so yum won't use it by default. Yum will **then** just ignore the repository **until** you permanently **enable** it again or use **--enablerepo** **for** temporary usage:

```
yum-config-manager --disable updates
```

4. Configure the failing repository to be skipped, **if** it is unavailable. Note that yum will try to contact the repo. when it runs most commands, so will have to try and fail each time (and thus. yum will be much slower). If it is a very temporary problem though, this is often a **nice compromise**:

```
yum-config-manager --save --setopt=updates.skip_if_unavailable=true
```

```
Cannot retrieve metalink for repository: updates/21/x86_64. Please verify its path and try
```

So this is a pretty scary error, but the actual cause is pretty tame. This happens once on almost every new Odroid install. The root cause is actually that the internal clock is set way back in the past to a time where the repository didn't exist. So when yum tries to check for the repository certificates, none of them are valid yet. To verify that this was the problem use the command:

```
$ date
```

If your Odroid is kickin' it in the '70s. You've found the problem. This is an easy fix. Run the following commands. The first updates the current time using the ntp server on campus. The second updates the hardware clock, which persists between power cycles provided you have that CMOS battery installed. If not, you'll have to do this every time.

```
$ ntpdate ntp.sdsmt.edu
$ hwclock -w
```

There was, last time I spoke to Scott, a potential problem with a dtb file not being updated in the kernel. If this happens and causes problems, just get a new image and repeat this process. Scott was hoping to have this fixed soon, so it's probably already done by now, but I'll leave this warning here anyway.

Bibliography

- [1] Paul Oliver, *Guppies - Evolving neural networks (w.i.p.)*, <https://youtu.be/tCPzYM7B338>, 2012.